

Project 1 Report

Team "RemCode":

Matvey Plevako, Mohamad Ziad AlKabakibi

Introduction

Topics, covered in the project:

- Understanding Docker-machine and Docker Swarm deployment
- Container Docker Cluster farm deployment
- Docker Container migration - Application Distribution
- Memory management
- Image size compression
- Building distributed Interactive web application

Application description:

RemCode is a website for remote code execution. Users can input code text, stdin to pass to the program and select language from the list. Then, they can view the results of their program.

Requirements for application:

- Web applications shouldn't execute any user code.
- Web applications should submit tasks to the task queue.
- Workers should connect to the task queue and pick tasks as they appear.
- Workers should execute the code with stdin passed with source code.
- After code execution, both successful and failed, workers should save results to the database.
- Users can view results of their program in a new window.
- Both Web application and Worker should be **scalable**.

Project

Understanding Docker-machine and Docker Swarm deployment

1. What is Docker-machine and what is it used for?

Docker-machine is a tool that is used for installing, setting and managing remote Docker hosts easily, moreover these hosts can be placed on different physical machines.

We can use docker-machine when we want to create a deployment environment for our application and manage all the micro-services running on it. With different drivers, docker-machine can be executed on different platforms for remote or local hosting.

2. What is Docker Swarm, what is it used for and why is it important in Containers Orchestration?

Docker Swarm is a system that is used for merging a set of Docker containers into a single coherent cluster and managing it. One of the key benefits associated with the operation of a docker swarm is the high level of availability offered for applications. In a docker swarm, there are typically several worker nodes and at least one manager node that is responsible for handling the worker nodes' resources efficiently and ensuring that the cluster operates efficiently. Moreover, it provides a convenient interface for managing containers on these hosts, services and distributing them among connected worker and manager nodes.

3. Install Docker-machine based on your virtualization platform (VirtualBox, Hyper-V, VMware), create a Machine (named Master), and collect some relevant information for you.

- docker-machine driver is virtual box
- Creating a new docker-machine host. we select virtualbox driver and name it "Master"

```
demo@myvm:~$ docker-machine create --driver virtualbox Master
Running pre-create checks...
Creating machine...
(Master) Copying /home/demo/.docker/machine/cache/boot2docker.iso to /home/demo/.docker/machine/machines/Master/boot2docker.iso...
(Master) Creating VirtualBox VM...
(Master) Creating SSH key...
(Master) Starting the VM...
(Master) Check network to re-create if needed...
(Master) Waiting for an IP...
Waiting for machine to be running, this may take a few minutes...
Detecting operating system of created instance...
Waiting for SSH to be available...
Detecting the provisioner...
Provisioning with boot2docker...
Copying certs to the local machine directory...
Copying certs to the remote machine...
Setting Docker configuration on the remote daemon...
Checking connection to Docker...
Docker is up and running!
To see how to connect your Docker Client to the Docker Engine running on this virtual machine, run: docker-machine env Master
```

- Running docker-machine ls after creating several docker-machines

```
demo@myvm:~$ docker-machine ls
```

NAME	ACTIVE	DRIVER	STATE	URL	SWARM	DOCKER	ERRORS
Master	-	virtualbox	Running	tcp://192.168.99.100:2376		v19.03.12	
Worker1	-	virtualbox	Running	tcp://192.168.99.101:2376		v19.03.12	
Worker2	-	virtualbox	Running	tcp://192.168.99.102:2376		v19.03.12	

Stopping and start docker-machine

```
demo@myvm:~$ docker-machine stop Worker2
Stopping "Worker2"...
Machine "Worker2" was stopped.
demo@myvm:~$ docker-machine ls
```

NAME	ACTIVE	DRIVER	STATE	URL	SWARM	DOCKER	ERRORS
Master	-	virtualbox	Running	tcp://192.168.99.100:2376		v19.03.12	
Worker1	-	virtualbox	Running	tcp://192.168.99.101:2376		v19.03.12	
Worker2	-	virtualbox	Stopped			Unknown	

```
demo@myvm:~$ docker-machine start Worker2
Starting "Worker2"...
(Worker2) Check network to re-create if needed...
(Worker2) Waiting for an IP...
Machine "Worker2" was started.
Waiting for SSH to be available...
Detecting the provisioner...
Started machines may have new IP addresses. You may need to re-run the `docker-machine env` command.
```

docker-machine rm

```
demo@myvm:~$ docker-machine rm Worker1
About to remove Worker1
WARNING: This action will delete both local reference and remote instance.
Are you sure? (y/n): y
Successfully removed Worker1
```

4. Create two Workers as well. Later we will connect them into one swarm. Make a screenshot for docker-machine ls command. You should have 3 running machines.

```
demo@myvm:~$ docker-machine ls
```

NAME	ACTIVE	DRIVER	STATE	URL	SWARM	DOCKER	ERRORS
Master	-	virtualbox	Running	tcp://192.168.99.100:2376		v19.03.12	
Worker1	-	virtualbox	Running	tcp://192.168.99.101:2376		v19.03.12	
Worker2	-	virtualbox	Running	tcp://192.168.99.102:2376		v19.03.12	

Container Docker Cluster farm deployment

5. Now that Docker Swarm is enabled, deploy a true container cluster farm across many Dockerized virtual machines. (One master and two workers). Verify the Docker Swarm status, identify the Master node(s), and how many workers active exist. Take as many screenshots as you need to explain the process.

Firstly, we need to initiate a Docker swarm with address and port to join.

```
docker@Master:~$ docker swarm init --advertise-addr 192.168.99.100:2377
Swarm initialized: current node (n6nqjh7tao9gxqo891gst18s9) is now a manager.

To add a worker to this swarm, run the following command:

    docker swarm join --token SWMTKN-1-1dy8cu0t4lqms592rnvovsxbxm7yrtbl62h416io5j6j41v-8g23nmzxtm268u3pnvfj003m0 192.168.99.100:2377

To add a manager to this swarm, run 'docker swarm join-token manager' and follow the instructions.
```

Then, we can check info about our swarm cluster

```

docker@Master:~$ docker info
Client:
 Debug Mode: false

Server:
 Containers: 0
  Running: 0
  Paused: 0
  Stopped: 0
 Images: 0
 Server Version: 19.03.12
 Storage Driver: overlay2
  Backing Filesystem: extfs
  Supports d_type: true
  Native Overlay Diff: true
 Logging Driver: json-file
 Cgroup Driver: cgroupfs
 Plugins:
  Volume: local
  Network: bridge host ipvlan macvlan null overlay
 Log: awslogs fluentd gcplogs gelf journald json-file local logentries splunk syslog
 Swarm: active
  NodeID: thj3ytc2o73pljwgio18adb4a
  Is Manager: true
  ClusterID: ac2mttba01k7kfkccqeldi7c57
 Managers: 1
 Nodes: 1
 Default Address Pool: 10.0.0.0/8
 SubnetSize: 24
 Data Path Port: 4789
 Orchestration:

```

Join worker nodes to swarm (commands for joining are too long, so they are compressed)

```

<dy8cu0t4lqgms592rnvocwvsbxqm7yrtbl62h416io5j6j41v-8g23nmzxtm268u3pnvfj003m0 192.168.99.100:2377
This node joined a swarm as a worker.
docker@Worker1:~$ exit
logout
demo@nyvm:~$ docker-machine ssh Worker2
( ' > ' )
 /) TC ( \   Core is distributed with ABSOLUTELY NO WARRANTY.
 (/ - _ - _ \)   www.tinycorelinux.net

<dy8cu0t4lqgms592rnvocwvsbxqm7yrtbl62h416io5j6j41v-8g23nmzxtm268u3pnvfj003m0 192.168.99.100:2377
This node joined a swarm as a worker.

```

list all active nodes in swarm

```

docker@Master:~$ docker node ls

```

ID	HOSTNAME	STATUS	AVAILABILITY	MANAGER STATUS	ENGINE VERSION
n6nqjh7tao9gxqo891gsti8s9 *	Master	Ready	Active	Leader	19.03.12
oek8fn09t7eayw037tqoeyz2z	Worker1	Ready	Active		19.03.12
y0mt9uulugck7s4cfu5nxhgdx	Worker2	Ready	Active		19.03.12

6. How can a Worker be promoted to Master and vice versa? Please explain if special requirements are needed to perform this action? Perform the process and explain it.

We can promote a worker node to the manager role using manager node. This is useful when a manager node becomes unavailable or if we want to take a manager offline for maintenance. Similarly, we can demote a manager node to the worker role.

Special requirements are we must always maintain a quorum of manager nodes in the swarm, in case of master node failures.

1. To promote worker node to the manager node:

```

docker@Master:~$ docker node ls

```

ID	HOSTNAME	STATUS	AVAILABILITY	MANAGER STATUS	ENGINE VERSION
n6nqjh7tao9gxqo891gsti8s9 *	Master	Ready	Active	Leader	19.03.12
oek8fn09t7eayw037tqoeyz2z	Worker1	Ready	Active		19.03.12
y0mt9uulugck7s4cfu5nxhgdx	Worker2	Ready	Active		19.03.12

```

docker@Master:~$ docker node promote Worker1
Node Worker1 promoted to a manager in the swarm.
docker@Master:~$ docker node ls

```

ID	HOSTNAME	STATUS	AVAILABILITY	MANAGER STATUS	ENGINE VERSION
n6nqjh7tao9gxqo891gsti8s9 *	Master	Ready	Active	Leader	19.03.12
oek8fn09t7eayw037tqoeyz2z	Worker1	Ready	Active	Reachable	19.03.12
y0mt9uulugck7s4cfu5nxhgdx	Worker2	Ready	Active		19.03.12

we promoted worker node and now we can perform cluster management from this node

2. To demote manager to worker:

```

docker@Master:~$ docker node ls

```

ID	HOSTNAME	STATUS	AVAILABILITY	MANAGER STATUS	ENGINE VERSION
n6nqjh7tao9gxqo891gsti8s9 *	Master	Ready	Active	Leader	19.03.12
oek8fn09t7eayw037tqoeyz2z	Worker1	Ready	Active	Reachable	19.03.12
y0mt9uulugck7s4cfu5nxhgdx	Worker2	Ready	Active		19.03.12

```

docker@Master:~$ docker node demote Worker1
Manager Worker1 demoted in the swarm.
docker@Master:~$ docker node ls

```

ID	HOSTNAME	STATUS	AVAILABILITY	MANAGER STATUS	ENGINE VERSION
n6nqjh7tao9gxqo891gsti8s9 *	Master	Ready	Active	Leader	19.03.12
oek8fn09t7eayw037tqoeyz2z	Worker1	Ready	Active		19.03.12
y0mt9uulugck7s4cfu5nxhgdx	Worker2	Ready	Active		19.03.12

7. Deploy a simple Web page, e.g Nginx, showing the hostname of the host node it is running upon, and validate that its instances are spreading across the servers previously deployed on your farm.

Solution: We created a simple html page and displayed \$hostname variable that is passed from nginx and served it with nginx server.

Created Dockerfile:

```
FROM nginx:alpine
COPY ./nginx.conf /etc/nginx/nginx.conf
COPY ./index.html /www/media/index.html
```

Created HTML page:

```
<html>
<body>

<!--# echo var="hostname" default="unknown_host" -->: <!--# echo var="server_port" default="unknown_port" -->

</body>
</html>
```

Created nginx.conf and enabled ssi to display \$hostname variable

```
http {
    server {
        listen 80 default_server;
        listen [::]:80 default_server;

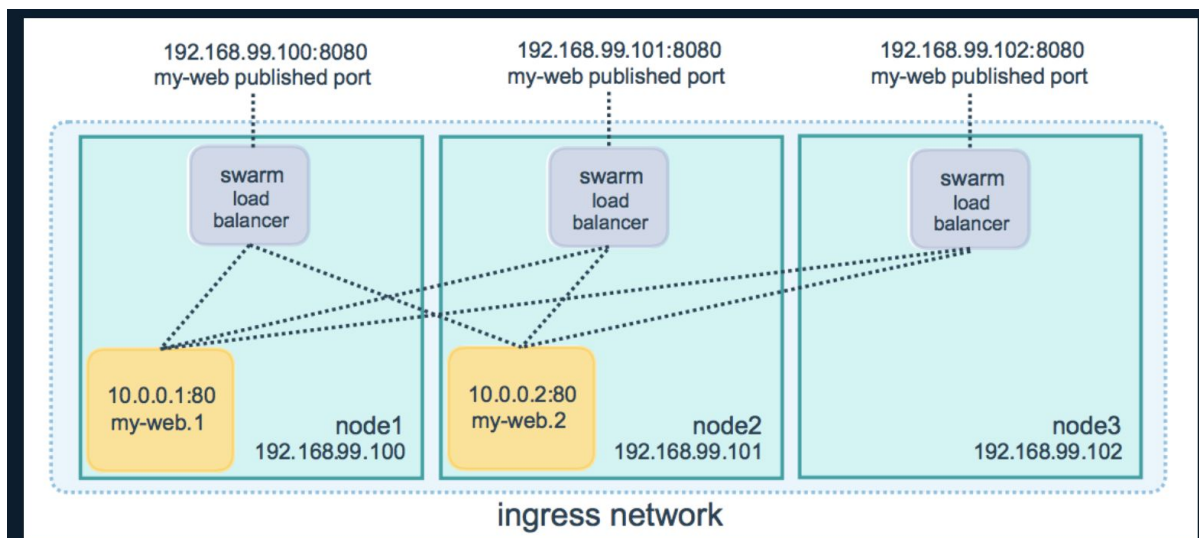
        location / {
            ssi on;
            root /www/media;
        }
    }
}
events { }
```

Spreading web-app for different hosts

```
docker@Master:~$ docker service create --name my-web --publish published=8080,target=80 --replicas 3 remcoded/nginx-page:latest
qrbfzke5il3rcbeec80wxcwzc
overall progress: 3 out of 3 tasks
1/3: running [=====]
2/3: running [=====]
3/3: running [=====]
verify: Service converged
docker@Master:~$ docker service ps my-web
```

ID	PORTS	NAME	IMAGE	NODE	DESIRED STATE	CURRENT STATE	ERROR
tnz670xxl1zk		my-web.1	remcoded/nginx-page:latest	Master	Running	Running 13 seconds ago	
vxphc1r5we15		my-web.2	remcoded/nginx-page:latest	Worker1	Running	Running 14 seconds ago	
sw2sgjgag1ca		my-web.3	remcoded/nginx-page:latest	Worker2	Running	Running 14 seconds ago	

When we publish a port from swarm the following happens: Each request will be processed by swarm load balancer and addressed to node according to the routing algorithm. Such a mesh provides a fault-tolerant system for the application.



Let's ping our service from master node.

```
docker@Master:~$ curl 0.0.0.0:8080
<html>
<body>

134bdb77eb15:80

</body>
</html>
docker@Master:~$ curl 0.0.0.0:8080
<html>
<body>

185c628e640e:80

</body>
</html>
docker@Master:~$ curl 0.0.0.0:8080
<html>
<body>

fb086571b64e:80

</body>
</html>
docker@Master:~$ curl 0.0.0.0:8080
<html>
<body>

134bdb77eb15:80

</body>
</html>
docker@Master:~$
```

As we can see, each time our request gets processed by a new node.

Let's scale down service to 2 nodes and check how requests will be distributed

```

docker@Master:~$ docker service scale my-web=2
my-web scaled to 2
overall progress: 2 out of 2 tasks
1/2: running [=====]
2/2: running [=====]
verify: Service converged
docker@Master:~$ curl 0.0.0.0:8080
<html>
<body>

185c628e640e:80

</body>
</html>
docker@Master:~$ curl 0.0.0.0:8080
<html>
<body>

fb086571b64e:80

</body>
</html>
docker@Master:~$ curl 0.0.0.0:8080
<html>
<body>

185c628e640e:80

</body>
</html>

```

Let's also test how requests are redistributed if we query each machine in swarm:

```

demo@myvm:~$ docker-machine ls
NAME      ACTIVE  DRIVER      STATE      URL                  SWARM      DOCKER      ERRORS
Master    -       virtualbox   Running    tcp://192.168.99.100:2376  v19.03.12
Worker1   -       virtualbox   Running    tcp://192.168.99.101:2376  v19.03.12
Worker2   -       virtualbox   Running    tcp://192.168.99.102:2376  v19.03.12
demo@myvm:~$ ^C
demo@myvm:~$ curl 192.168.99.100:8080
<html>
<body>

d33bdf505278:80

</body>
</html>
demo@myvm:~$ curl 192.168.99.101:8080
<html>
<body>

d33bdf505278:80

</body>
</html>
demo@myvm:~$ curl 192.168.99.102:8080
<html>
<body>

30fcb4f0cd2c:80

</body>
</html>

```

As we can see, each machine can be queried and we will get a response even if there are only 2 instances running these services.

8. How to scale instances in the Docker Swarm? Could it be done automatically?

We can scale services manually using

`docker service scale <service-name>=<number-of-instances>`

```

docker@Master:~$ docker service scale my-web=2
my-web scaled to 2
overall progress: 2 out of 2 tasks
1/2: running [=====]
2/2: running [=====]
verify: Service converged

```

There are no built-in functions in docker swarm for automated scaling, because hosts are created with docker-machine with different drivers. However, it can be automated by scripts using metrics.

Docker Container migration - Application Distribution

9. Validate that when a node goes down a new instance is launched. Show how the redistribution of the instances can happen when the dead node comes back alive.

Initial state:

```
docker@Master:~$ docker node ls
```

ID	HOSTNAME	STATUS	AVAILABILITY	MANAGER STATUS	ENGINE VERSION
n6nqjh7tao9gxqo891gsti8s9 *	Master	Ready	Active	Leader	19.03.12
oek8fn09t7eayw037tqoeyz2z	Worker1	Ready	Active		19.03.12
y0mt9uulugck7s4cfu5nxhgdx	Worker2	Ready	Active		19.03.12


```
docker@Master:~$ docker service ps my-web
```

ID	PORTS	NAME	IMAGE	NODE	DESIRED STATE	CURRENT STATE	ERROR
bvtma8naaite		my-web.1	remcoded/nginx-page:latest	Worker2	Running	Running 21 seconds ago	
w4deoptdbzai		my-web.2	remcoded/nginx-page:latest	Master	Running	Running 21 seconds ago	
rc68a0h808hm		my-web.3	remcoded/nginx-page:latest	Worker1	Running	Running 21 seconds ago	

Next, we will stop Worker2 and see distribution of instances

```
demo@myvm:~$ docker-machine stop Worker2
Stopping "Worker2"...
Machine "Worker2" was stopped.
demo@myvm:~$ docker-machine ls
```

NAME	ACTIVE	DRIVER	STATE	URL	SWARM	DOCKER	ERRORS
Master	-	virtualbox	Running	tcp://192.168.99.100:2376		v19.03.12	
Worker1	-	virtualbox	Running	tcp://192.168.99.101:2376		v19.03.12	
Worker2	-	virtualbox	Stopped			Unknown	


```
docker@Master:~$ docker service ps my-web
```

ID	PORTS	NAME	IMAGE	NODE	DESIRED STATE	CURRENT STATE	ERROR
9z9o8d5ze6zv		my-web.1	remcoded/nginx-page:latest	Master	Running	Running 25 seconds ago	
bvtma8naaite		_ my-web.1	remcoded/nginx-page:latest	Worker2	Shutdown	Running 2 minutes ago	
w4deoptdbzai		my-web.2	remcoded/nginx-page:latest	Master	Running	Running 2 minutes ago	
rc68a0h808hm		my-web.3	remcoded/nginx-page:latest	Worker1	Running	Running 2 minutes ago	

As we can see, master node increases the number of instances it runs.

Next, we will start Worker2 again.

```
demo@myvm:~$ docker-machine start Worker2
Starting "Worker2"...
(Worker2) Check network to re-create if needed...
(Worker2) Waiting for an IP...
Machine "Worker2" was started.
Waiting for SSH to be available...
Detecting the provisioner...
Started machines may have new IP addresses. You may need to re-run the `docker-machine env` command.
```

According to docker documentation, when node goes up, automatic redistribution does not happen in order to prevent interrupting running services

Force the swarm to rebalance

Generally, you do not need to force the swarm to rebalance its tasks. When you add a new node to a swarm, or a node reconnects to the swarm after a period of unavailability, the swarm does not automatically give a workload to the idle node. This is a design decision. If the swarm periodically shifted tasks to different nodes for the sake of balance, the clients using those tasks would be disrupted. The goal is to avoid disrupting running services for the sake of balance across the swarm. When new tasks start, or when a node with running tasks becomes unavailable, those tasks are given to less busy nodes. The goal is eventual balance, with minimal disruption to the end user.

In Docker 1.13 and higher, you can use the `--force` or `-f` flag with the `docker service update` command to force the service to redistribute its tasks across the available worker nodes. This causes the service tasks to restart. Client applications may be disrupted. If you have configured it, your service uses a rolling update.

Automatic task distribution is not happening, so to have the same number of containers on each host we have to force rebalance nodes.

```

docker@Master:~$ docker service update --force my-web
my-web
overall progress: 3 out of 3 tasks
1/3: running [=====>]
2/3: running [=====>]
3/3: running [=====>]
verify: Service converged
docker@Master:~$ docker service ps my-web

```

ID	NAME	IMAGE	NODE	DESIRED STATE	CURRENT STATE	ERROR
h9ggp00n952i	my-web.1	remcoded/nginx-page:latest	Worker2	Running	Running 3 minutes ago	
9z9o8d5ze6zv	_ my-web.1	remcoded/nginx-page:latest	Master	Shutdown	Shutdown 3 minutes ago	
bvtma8naalte	_ my-web.1	remcoded/nginx-page:latest	Worker2	Shutdown	Shutdown 8 minutes ago	
n9bi6vnwlpfq	my-web.2	remcoded/nginx-page:latest	Master	Running	Running 3 minutes ago	
w4deoptdbzai	_ my-web.2	remcoded/nginx-page:latest	Master	Shutdown	Shutdown 3 minutes ago	
l5rprhyv8s2j	my-web.3	remcoded/nginx-page:latest	Worker1	Running	Running 3 minutes ago	
rc68a0h808hm	_ my-web.3	remcoded/nginx-page:latest	Worker1	Shutdown	Shutdown 3 minutes ago	

As we can see, container was taken down on Master and was created on Worker 2

10. Perform some update in your application, a minor change in your sample application for example. How to replicate the changes in the rest of the farm servers?

We need to tell docker to update all nodes that don't correspond to, for example, the latest image of service. The update will happen as follows:

1. Update will happen sequentially, so service will keep running on other nodes, while updating other
2. Task will be shutted down on running node
3. Task will be recreated with new image and updater will proceed to other nodes
4. After all nodes are updated, the updater will verify that they are running without error for at least 5 seconds and finish.

```

docker@Master:~$ docker service update --image remcoded/nginx-page:latest my-web
my-web
overall progress: 2 out of 2 tasks
1/2: running [=====>]
2/2: running [=====>]
verify: Service converged

```

11. It is a good practice to monitor performance and logs on your servers farm. How can this be done with Docker Swarm? Could it be just CLI or maybe GUI?

There are good GUI tools for visualizing performance such as <https://docs.docker.com/config/daemon/prometheus/> can be used.

We create this service on already running nodes and it monitors performance metrics of Docker containers running on different hosts.

Playing with Memory

12. Please explain what is “Out Of Memory Exception (OOMF)”, how it could affect Docker services, and which configuration can be set to avoid this issue?

We can refer to the official Docker documentation. OOME is dangerous because it can make the system unstable and kill other processes.

Understand the risks of running out of memory

It is important not to allow a running container to consume too much of the host machine's memory. On Linux hosts, if the kernel detects that there is not enough memory to perform important system functions, it throws an `OOMF`, or `Out Of Memory Exception`, and starts killing processes to free up memory. Any process is subject to killing, including Docker and other important applications. This can effectively bring the entire system down if the wrong process is killed.

Most obvious solution is we can prevent OOME by limiting use of memory by containers, but we also have to follow these rules to prevent it:

You can mitigate the risk of system instability due to OOME by:

- Perform tests to understand the memory requirements of your application before placing it into production.
- Ensure that your application runs only on hosts with adequate resources.
- Limit the amount of memory your container can use, as described below.
- Be mindful when configuring swap on your Docker hosts. Swap is slower and less performant than memory but can provide a buffer against running out of system memory.
- Consider converting your container to a service, and using service-level constraints and node labels to ensure that the application runs only on hosts with enough memory

We can enforce Docker to limit memory usage with the following commands:

Limit a container's access to memory

Docker can enforce hard memory limits, which allow the container to use no more than a given amount of user or system memory, or soft limits, which allow the container to use as much memory as it needs unless certain conditions are met, such as when the kernel detects low memory or contention on the host machine. Some of these options have different effects when used alone or when more than one option is set.

Most of these options take a positive integer, followed by a suffix of `b`, `k`, `m`, `g`, to indicate bytes, kilobytes, megabytes, or gigabytes.

Option	Description
<code>-m</code> or <code>--memory=</code>	The maximum amount of memory the container can use. If you set this option, the minimum allowed value is <code>4m</code> (4 megabyte).
<code>--memory-swap *</code>	The amount of memory this container is allowed to swap to disk. See <code>--memory-swap</code> details.
<code>--memory-swappiness</code>	By default, the host kernel can swap out a percentage of anonymous pages used by a container. You can set <code>--memory-swappiness</code> to a value between 0 and 100, to tune this percentage. See <code>--memory-swappiness</code> details.
<code>--memory-reservation</code>	Allows you to specify a soft limit smaller than <code>--memory</code> which is activated when Docker detects contention or low memory on the host machine. If you use <code>--memory-reservation</code> , it must be set lower than <code>--memory</code> for it to take precedence. Because it is a soft limit, it does not guarantee that the container doesn't exceed the limit.
<code>--kernel-memory</code>	The maximum amount of kernel memory the container can use. The minimum allowed value is <code>4m</code> . Because kernel memory cannot be swapped out, a container which is starved of kernel memory may block host machine resources, which can have side effects on the host machine and on other containers. See <code>--kernel-memory</code> details.
<code>--oom-kill-disable</code>	By default, if an out-of-memory (OOM) error occurs, the kernel kills processes in a container. To change this behavior, use the <code>--oom-kill-disable</code> option. Only disable the OOM killer on containers where you have also set the <code>-m/--memory</code> option. If the <code>-m</code> flag is not set, the host can run out of memory and the kernel may need to kill the host system's processes to free memory.

13. Deploy a docker container with at least 15% of CPU every second for memory efficiency.

From the Docker documentation, we can use `cpus` parameter to achieve this.

`--cpus=<value>`

Specify how much of the available CPU resources a container can use. For instance, if the host machine has two CPUs and you set `--cpus="1.5"`, the container is guaranteed at most one and a half of the CPUs. This is the equivalent of setting `--cpu-period="100000"` and `--cpu-quota="150000"`. Available in Docker 1.13 and higher.

To deploy my-web and limit it to 15% usage:

```
demo@myvm:~$ docker run --cpus="0.15" my-web
```

Compression

14. Verify the size of the Docker images that you're working with. Can this size be reduced and how can we achieve this?

Verify the size of my-web:

```
docker@Master:~$ docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
remcodeds/nginx-page	latest	b42cb789c1ff	18 hours ago	133MB
remcodeds/nginx-page	<none>	85683af5f261	18 hours ago	133MB
remcodeds/nginx-page	<none>	85683af5f261	18 hours ago	133MB

There are some common advices for reducing image size:

1. Change base image to alpine.
2. (e.g) FROM nginx:alpine
3. Use multi-stage builds to minimize unused artifacts
4. Don't install unnecessary for production debug tools
5. Reduce number of layers and perform operations simultaneously
6. run installs with --no-install-recommends to install only main dependencies

After applying these suggestions to my-web image:

```
docker@Master:~/step7$ docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
remcodeds/nginx-page	latest	5838c1c1f978	7 seconds ago	22.1MB
remcodeds/nginx-page	<none>	b42cb789c1ff	19 hours ago	133MB
remcodeds/nginx-page	<none>	85683af5f261	18 hours ago	133MB

So we reduced size from 133MB to 22.1MB

Web application

Cluster is currently Up and running on <http://34.77.110.39:8888/>

First, let's start with creating python scripts for Web-application and Worker.
Web-application is written with Django.

We will present here only important parts of code.

view.py

```
from django.http import HttpResponseRedirect, HttpResponse
from django.shortcuts import render
import hashlib
import time

from redis import from_url

from .forms import GetCode

def get_data(request):
    r = from_url("redis://redis:6379")
    # if this is a POST request we need to process the form data
    if request.method == 'POST':
        # create a form instance and populate it with data from the request:
        form = GetCode(request.POST)
        # check whether it's valid:
        if form.is_valid():
            code = form.cleaned_data['code']
            stdin = form.cleaned_data['stdin']
            lang = form.cleaned_data['language']
            token = hashlib.sha256(code.encode()).hexdigest()
            r.lpush("tasks", f"{token}\n{lang}\n{stdin}{token}{code}")
            url = f'/token/{token}/'
            time.sleep(1)
            return HttpResponseRedirect(url)

    # if a GET (or any other method) we'll create a blank form
    else:
        form = GetCode()

    return render(request, 'code_submit.html', {'form': form})

def get_code_result(request, token):
    r = from_url("redis://redis:6379")
    res = r.get(token)
    if res:
        return HttpResponse(res.decode())
    else:
        return render(request, 'wait_for_code.html')
```

Important points:

1. get_data process data, that user inputed and redirects user to the page with result
2. get_code_results returns program output after execution.

worker.py

```
from redis import from_url
import subprocess
import sys

import logging

def run(cmd, stdin=""):
    proc = subprocess.Popen(cmd,
                             stdin=subprocess.PIPE,
                             stdout=subprocess.PIPE,
                             stderr=subprocess.PIPE,
                             shell=True
                             )

    stdout, stderr = proc.communicate(input=stdin.encode())

    return proc.returncode, stdout, stderr

def execute_code(code_text, stdin, language, token):
    r = from_url("redis://redis:6379")
    if language == "python3":
        exit_code, out, err = run_py(code_text, stdin)
    else:
        exit_code, out, err = run_gcc(code_text, stdin)

    logging.info(out)
    logging.info(err)
    logging.info(exit_code)
    print("out: {}".format(out))
    print("err: {}".format(err))
    print("exit: {}".format(exit_code))
    if exit_code == 0:
        if out:
            r.set(token, out)
        else:
            r.set(token, "no output was produced by code")
    else:
        r.set(token, err)

def run_py(code_text, stdin):
    with open("tmp.py", "w") as file:
        file.write(code_text)

    exit_code, out, err = map(lambda x: x.decode() if type(x) == bytes else x,
                              run([f"{sys.executable} tmp.py"], stdin))
    run(["rm tmp.py"])

    return exit_code, out, err

def run_gcc(code_text, stdin):
    with open("tmp.c", "w") as file:
        file.write(code_text)
```



```

exit_code, out, err = map(lambda x: x.decode() if type(x) == bytes else x,
                           run(["gcc tmp.c && ./a.out"], stdin))
run(["rm tmp.c"])
run(["rm a.out"])

return exit_code, out, err

def main():
    while True:
        r = from_url("redis://redis:6379")
        token, lang_input_code = r.brpop("tasks")[1].decode().split("\n", 1)
        lang, input_code = lang_input_code.split("\n", 1)
        stdin, code = input_code.split(token, 1)
        execute_code(code, stdin, lang, token)

if __name__ == '__main__':
    main()

```

Important points:

1. program blocks on redis list to avoid pulling
2. After token is received, it is processed and passed to execute_code
3. execute_code depending on the language, execute the source program
4. result is written into the redis database

Dockerfile for Web-App

```

FROM python:3
ENV PYTHONUNBUFFERED 1
RUN mkdir /code
WORKDIR /code
COPY requirements.txt /code/
RUN pip install -r requirements.txt
COPY . /code/

```

Dockerfile for Worker

```

FROM python:3
ENV DEBIAN_FRONTEND noninteractive
RUN apt-get update && apt-get -y install gcc mono-mcs && rm -rf /var/lib/apt/lists/*
RUN pip install redis
COPY worker.py .
CMD python worker.py

```

docker-compose.yml

```

version: '3'

services:
  postgres:
    image: postgres:alpine
    environment:
      - POSTGRES_DB=postgres
      - POSTGRES_USER=postgres
      - POSTGRES_PASSWORD=postgres

  redis:

```

```

image: redis:alpine

webapp:
  build: remCodeWeb
  command: python manage.py runserver 0.0.0.0:8000
  ports:
    - "8000:8000"
  depends_on:
    - postgres
    - redis

worker:
  build: remCodeWorker
  depends_on:
    - redis

```

it will be used with docker stack deploy

Deployment:

1. push all locally build images to the repository
 - a. docker push remcodeds/worker
 - b. docker push remcodeds/web-app
2. replace all locally build images in docker-compose.yml with pushed ones
3. deploy stack

```

docker@Master:~/step7$ docker stack deploy --compose-file docker-compose.yml remcode
Creating network remcode_default
Creating service remcode_webapp
Creating service remcode_worker
Creating service remcode_postgres
Creating service remcode_redis

```

4. Check status of stack

```

docker@Master:~/step7$ docker stack services remcode

```

ID	NAME	MODE	REPLICAS	IMAGE	PORTS
ex9nxwslrl32	remcode_redis	replicated	1/1	redis:alpine	
hys5jvxpqeh9	remcode_postgres	replicated	1/1	postgres:alpine	
npw0v9w15voz	remcode_worker	replicated	1/1	remcodeds/worker:latest	
oas4j56a3bhb	remcode_webapp	replicated	1/1	remcodeds/web-app:latest	*:8000->8000/tcp

as we can see, every service has been replicated=1 and web-app has opened port

5. As most of the workload will be on workers, we can scale them to 4 instances.

```

docker@Master:~/step7$ docker service scale remcode_worker=4
remcode_worker scaled to 4
overall progress: 4 out of 4 tasks
1/4: running [=====>]
2/4: running [=====>]
3/4: running [=====>]
4/4: running [=====>]
verify: Service converged

```

6. We can also scale web-app to 2 instances.

```

docker@Master:~/step7$ docker service scale remcode_webapp=2
remcode_webapp scaled to 2
overall progress: 2 out of 2 tasks
1/2: running [=====>]
2/2: running [=====>]
verify: Service converged

```

7. The final picture of deployed services looks like this:

```

docker@Master:~/step7$ docker stack services remcode

```

ID	NAME	MODE	REPLICAS	IMAGE	PORTS
a901d3xsttrx	remcode_postgres	replicated	1/1	postgres:alpine	
hjd5a0pp3c2	remcode_worker	replicated	4/4	remcodeds/worker:latest	
uicd7a31h202	remcode_redis	replicated	1/1	redis:alpine	
wgqehcspqtex	remcode_webapp	replicated	2/2	remcodeds/web-app:latest	*:8000->8000/tcp

Web interface

The main page:

RemCode

Code

```
#include <stdio.h>
int main() {
    int number;

    printf("Enter an integer: ");

    // reads and stores input
    scanf("%d", &number);

    // displays output
    printf("You entered: %d", number);

    return 0;
}
```

Input for program

25

Language

c

Submit

Page with results of program:

RemCode


Result of your program:

Enter an integer: You entered: 25

Page with message if results are not ready yet:

RemCode

Your code is still executing please wait



Conclusion

What we learned

1. What is and how to use docker-machine
2. How to deploy our own swarm cluster
3. How to manage nodes in swarm and deploy new Manager nodes
4. How to build scalable applications
5. How to deploy applications in docker cluster and how to scale them
6. How to limit resources used by containerized application

Difficulties

1. Come up with scalable architecture for application
2. How to pass data between containers
3. How to build scalable worker application